




ViT rollout uses dense attention-only state after layer 0. Qwen3-VL is true visual forward, per temporal slice. Proxy is row-renormalized masked attention, not a trained Monarch factorization.