

Proxy overhead only: not end-to-end GPU latency

